

06 | Jan | 2023.


## Hidden Markov Models

# Markov Model

- Let's say there is a process which generates a sequence:

- Sequence of written words     *text*
- Sequence of spoken words (speech)
- DNA, RNA protein's amino acid sequence, etc.

- The observed sequence is governed by some context which itself keeps updating as the observations are generated.



$w_1, w_2, \dots, w_n$

$p(\text{obs} \mid \text{past history})$

- The probability of a sequence can be written as

$$P(w_{1:n}) = P(w_1) P(w_2 \mid w_1) P(w_3 \mid w_{1:2}) \dots P(w_n \mid w_{1:n-1})$$

past history

entire sequence with indices 1 to n

$w_1, w_2, w_3, \dots, w_{n-1}, w_n$

$w_{1:n}$

$w_{1:n}$

- This is the chain rule of probability.

- Modeling the sequence of words  $P(w_{1:n})$  requires learning each of the component terms.

$$P(w_3 \mid w_1, w_2)$$

$w_1$	$w_2$	$P(w_3)$
—	—	—
—	—	—

- We can do something smarter by modelling the context in a better way.

$$P(w_{1:n})$$

We want to model.

CT

CPT<sub>3</sub>

$P(w_1)$   $P(w_2 \mid w_1)$  needs to be learned

$P(w_1 w_2 \dots w_n)$  parameterized CPT

Valid sentence  $w_1 \dots w_n$   
Correct

Prob high.

↓  
has errors

Prob ↓.

Language Model

$w_{1,n}$  →

Conforms to a language  
well formed sentence

- The context sequence can be

- Observable

- Hidden

HMMs

Context  $\longrightarrow$  observation

$$P(w_k | \underbrace{w_{1:k-1}}_{\text{context}})$$

- An example of observable context model is the n-gram model of text.

- Limiting the context
- n-Gram model: Only the previous n-1 words have any effect on the probabilities of the next word.

trigram  $P(w_3 | w_{1,2})$

$P(w_5 | w_{3,4})$   $\frac{w_1}{\underbrace{\quad}_{n-1 \text{ previous words}}}$

unigram  $P(w_i)$   
 Bigram  $P(w_2 | w_1)$   
 $P(w_1 | w_0)$

- For a trigram model:

$$p(w_n | w_1, \dots, w_{n-1}) = p(w_n | w_{n-2}, w_{n-1})$$

past history
last two words

Due to limited context (assumption)

- Applying the trigram model gives:

LHS  
Joint  
distribution  
of  $w_{1:n}$

$$P(w_{1:n}) = P(w_1) P(w_2 | w_1) p(w_3 | w_{1,2}) \dots P(w_n | w_{n-2}, w_{n-1})$$

1      2      3 . . . . . 3 variables
CPTs

$$= P(w_1) P(w_2 | w_1) \prod_{i=3}^n P(w_i | w_{i-2}, w_{i-1})$$

$w_0$   
 $w_{-1}$

$$\prod_{i=1}^n P(w_i | w_{i-2}, w_{i-1})$$

Corpus.

- To create such a trigram model, we require the probability of every possible valid trigram.

$$\rightarrow \frac{C(w_{i-2} w_{i-1} w_i)}{C(w_{i-2} w_{i-1})}$$

- The estimate of the conditional probability

Training dataset

$e$ : estimate

$$P_e(w_i | w_{i-2}, w_{i-1}) = \frac{C(w_{i-2}, w_{i-1}, w_i)}{C(w_{i-2}, w_{i-1})}$$

Count


$$\left. \begin{matrix} w_1 & w_2 \\ \vdots & \vdots \end{matrix} \right| \begin{matrix} P(w_3) \\ 0 \\ 0 \\ 0 \end{matrix} \quad P(w_3 = \dots)$$

- "To create such a model we simply go through some training text"

Corpus

$$P_e('model' | 'such', 'a') = \frac{C('such', 'a', 'model')}{C('such', 'a')} = \frac{1}{1}$$

- After all the hard work, we now have a model which can generate the probability of any string.  $P(\underline{w_{1,n}})$
- This model is called a Markov chain.  $\rightarrow$  Context is observable.
- It can be represented as a graph of nodes and arcs
- Nodes are the states  $\rightarrow$  context
- Arcs represent transitions from one state to another.
  - An arc ~~e~~ is labelled with the observation (emitted symbol) and the probability of the observed symbol given the state





- Why is the state obvious here?


- Given a Markov chain, the probability of an observed sequence can be computed as the product of the probabilities of each transition in the path.

$$P(w_{1:n})$$

a b c d

$$P(s^1 | s^0, a) = \frac{P(s^1 | s^0, a)}{P(a | s^0)}$$

$$P(abcd) =$$





$$P(s^0) P(s^1 | s^0, a) P(s^2 | s^1, b) P(s^3 | s^2, c) P(s^4 | s^3, d)$$

## What if the states are hidden?


- Now consider the scenario when the states are non-obvious.
- That means, we are uncertain about the state sequence that generated the given observation sequence.
- That means,
  - there can be multiple state sequences that can lead to a given observation sequence.
  - a given state sequence can generate multiple observation sequences
  - for a given state, we can make transition to multiple possible states, even while generating the same output.

## Hidden Markov Models

- So now, transition to multiple states is possible from a given state, even while emitting the same observation (symbol).  
*// Difference from a Markov chain.*
- HMMs are a generalisation of Markov Chains in which a given state may have several transitions out of it, all with the same label (symbol).  
*and*
- Recall that in a Markov chain, given a state, the next state is certain (fixed) once you observe a given symbol.



- Since more than one successor state may have the same output, in general there may be several paths through an HMM that produce the same output *sequence*.



- In such cases, the probability of the output is the sum of the probabilities of all possible paths.
- The Markov chain cannot be used now.
  - What we need is a state transition diagram to represent an HMM.

## Moving from a Markov Chain to an HMM

- Consider that we are learning a trigram model.
- If the training corpus is sparse, it may happen that there is a trigram in the next text that never appeared in the training corpus.
  - In that case, a 0 probability will be assigned to the trigram, i.e. probability of the third word, given the first two words.
- One solution to this problem is to smooth the probabilities by also using the bigram and unigram probabilities.

$$P(w_3 | w_1 w_2) \quad \text{CPT}$$

$$P(w_n | w_{n-2}, w_{n-1}) = \lambda_1 P_e(w_n) + \lambda_2 P_e(w_n | w_{n-1}) + \lambda_3 P_e(w_n | w_{n-2}, w_{n-1})$$

Diagram illustrating the smoothing equation with handwritten annotations:

- The left side is labeled trigram.
- The first term is labeled uni (unigram) and Path 1 with a value of 0.1.
- The second term is labeled Bi (bigram) and Path 2 with a value of 0.3.
- The third term is labeled Trigram and Path 3 with a value of 0.6.
- Handwritten notes include:  $\lambda_1 + \lambda_2 + \lambda_3 = 1$  and  $\lambda_3 > \lambda_2 > \lambda_1$ .
- Small values 0.6, 0.3, and 0.1 are written near the terms.

- If we have missing trigram then the bigram and unigram probabilities take over.
- If we have missing trigram and bigram, then we fall back on the unigram probability.
- The probability of a transition has become an addition of multiple components (arcs) between the two states.
  - An HMM will allow multiple paths between two states.

## HMM state transition diagram for a trigram model

- Consider that we have seen the symbol ab.
- The next symbol in the sequence is a.
- So the next state is ba
- The trigram  $P(a|ab)$  itself can be represented as an HMM

$$P(a|ab) = {}_1P_e(a) + {}_2P_e(a|b) + {}_3P_e(a|ab)$$



- An HMM is a 4 tuple

$$s^1, S, W, E$$

- $S$ : the set of states
- $s^1$      $S$  is the initial state of the model
- $W$  : the set of output symbols generated/emitted/accepted
- $E$  : the set of edges or transitions.

- For each of the sets  $S, W, E$  we assume a canonical ordering of the elements

$$S = s^1, s^2, \dots, s$$

$$W = w^1, w^2, \dots, w$$

$$E = e^1, e^2, \dots, e$$

- The starting state of the HMM is the first element in the ordering of states.
- A transition is a 4-tuple  $s^i, s^j, w^k, p$  where  $s^i$  is the state from which the transition starts,  $s^j$  is the state where the transition ends,  $w^k$  is the output symbol generated by the model,  $p$  is the probability of taking the transition

- $p$  is the probability of the transition  $s^i \xrightarrow{w^k} s^j$
- No two transitions can have the same starting and ending states as well as the same output value.
- We can leave out the zero probability paths from the graphical representation of the HMM.
- Note that a state  $s$  can be the starting state for several transitions that have the same output symbol but go to different ending states.

- It is not possible to know what state the machine has gone into simply by looking at the output.
- Thus, the state sequence followed by an HMM is not deductible from the input. It is hidden.
- There are  $n+1$  states for  $n$  outputs

- The probability  $p$  associated with a transition  $s^i \xrightarrow{w_k} s^j$

$$\begin{aligned} P(s^i \xrightarrow{w_k} s^j) &= P(S_{t+1} = s^j, W_t = w^k | S_t = s^i) \\ &= P(s^j, w^k | s^i) \end{aligned}$$

- When writing  $P(s^j, w^k | s^i)$ , it is understood that  $s^i$  is the prior state and  $s^j$  is the next state.

- Markov models assume that the only information affecting the probability of an output, or of the next state, is the prior state.
- As per this Markov Assumption

$$P(w_n, s_{n+1} | w_{1,n-1}, s_{1,n}) = P(w_n, s_{n+1} | s_n) = P(s_n^i \xrightarrow{w_n^k} s_{n+1}^j)$$

- The probability of a sequence  $w_{1,n}$  is the probability of all possible paths through the HMM that can produce this sequence.

In other words:

$$P(w_{1,n}) = \sum_{s_{1,n+1}} P(w_{1,n}, s_{1,n+1})$$

## The Markov Assumption helps

$$P(w_{1:n}) = \prod_{s_{1:n+1}} P(w_{1:n}, s_{1:n+1})$$

$$P(w_{1:n}) = \prod_{s_{1:n+1}} P(s_1)P(w_1, s_2 | s_1)P(w_2, s_3 | w_1, s_{1,2}) \dots P(w_n, s_{n+1} | w_{1:n-1}, s_{1,n})$$

- The Markov Assumption helps here

$$P(w_{1:n}) = \prod_{s_{1:n+1}} \prod_{i=1}^n P(w_i, s_{i+1} | s_i)$$

Exercise: Graphically visualize the simplification offered by Markov Assumption



## An application of HMM: Part of Speech Tagging

- English words can be in more than one PoS class.
- PoS tags can be assigned using HMM.
- But to use HMM we need to phrase the problem of assigning PoS tags to the words as one of assigning probabilities to the input text.
- We assume that the HMM generates outputs which are the words of the corpus.
- We assume that there is some connection between the states and the transitions of the tags.

- Earlier we had for Language Model (LM)

$$P(w_{1,n}) = \prod_{s_{1,n+1}} P(w_{1,n}, s_{1,n+1})$$

- For PoS tagging, we assume correspondence between states and the tags.
- So,

$$P(w_{1,n}) = \prod_{t_{1,n+1}} P(w_{1,n}, t_{1,n+1})$$

- Here,  $t_{1,n+1}$  is a sequence of  $n + 1$  parts of speech or tags.

- The last tag  $t_{n+1}$  is the tag achieved after emitting  $w_n$ .
- So  $t_{n+1}$  is a tag which corresponds to the non-existent word  $w_{n+1}$
- We define the problem of PoS tagging as finding

$$\arg \max_{t_{1,n}} P(t_{1,n} | w_{1,n}) = \arg \max_{t_{1,n}} \frac{P(w_{1,n}, t_{1,n})}{P(w_{1,n})} = \arg \max_{t_{1,n}} P(w_{1,n}, t_{1,n})$$

